

Script for the program 'CLP_vs_Copper.py' in Python

Licensed under the Apache License, Version 2.0


<http://www.apache.org/licenses/LICENSE-2.0>

The program contains a study (linear regression + statistical interpretation) revealing how the Chilean peso has decoupled from the price of its main export, copper, after the Chilean civil unrest of October 2019

```
In [1]: # Import Python Libraries (pandas for data, matplotlib.pyplot for data plots,  
# matplotlib.image for image plots and statsmodels.formula.api for statistics)  
  
import pandas as pd  
import matplotlib.pyplot as plt  
import matplotlib.image as img  
import statsmodels.formula.api as smf
```

First we observe from a graph of bloomberg.com that the Chilean peso seems to have decoupled from copper prices after the Chilean civil unrest of October 2019:

```
In [3]: chart = img.imread('Bloomberg.png')  
plt.figure(figsize=(16, 16))  
plt.imshow(chart)  
plt.axis('off')  
plt.show()
```



From investing.com website:

<https://www.investing.com/currencies/usd-clp-historical-data>

<https://www.investing.com/commodities/copper-historical-data>

one can download csv files which contain the data of the currency pair USD-CLP and the copper price before the civil unrest, during the time frame Oct 06, 2017 - Oct 04, 2019.

```
In [4]: # Store the csv files in dataframes pesos_before and copper_before, and order them by ascending date  
  
pesos_before = pd.read_csv('CLP_before.csv')  
copper_before = pd.read_csv('Copper_before.csv')  
pesos_before, copper_before = pesos_before[::-1], copper_before[::-1]
```

```
In [5]: # Merge the data into one dataframe 'df_before', with 'Date' as (inner) merge criterion

df_before = pd.merge(left=pesos_before, right=copper_before, left_on='Date', right_on='Date')
df_before = [df_before['Date'], df_before['Price_x'], df_before['Price_y']]
headers = ['Date', 'USD_CLP', 'Copper_price']
df_before = pd.concat(df_before, axis=1, keys=headers)
```

Now we perform an ordinary least squares regression (OLS) between 'USD-CLP' and 'Copper price'. First an OLS object is generated, and then the fit() method is applied for fitting the regression line to the data. See more details on:

<https://www.geeksforgeeks.org/ordinary-least-squares-ols-using-statsmodels/>

https://en.wikipedia.org/wiki/Ordinary_least_squares

```
In [6]: # The attribute 'params' returns the list of estimated parameters for the model.
# These are the intercept and the slope of the OLS regression, params[0] and params[1], respectively.

fit_before = smf.ols(formula='USD_CLP~Copper_price', data=df_before).fit()
intercept_before, slope_before = fit_before.params[0], fit_before.params[1]
```

```
In [7]: # Equation of OLS regression (with column 'Best_fit'), and error between the 'Best_fit' and the obser


df_before['Best_fit'] = slope_before*df_before['Copper_price'] + intercept_before
df_before['Error'] = df_before['USD_CLP'] - df_before['Best_fit']
error_before = round((df_before['Error']**2).sum(), 1)
print(df_before.head())
```

| | Date | USD_CLP | Copper_price | Best_fit | Error |
|---|--------------|---------|--------------|------------|-----------|
| 0 | Oct 06, 2017 | 633.22 | 3.1135 | 636.494462 | -3.274462 |
| 1 | Oct 09, 2017 | 633.40 | 3.1170 | 635.989263 | -2.589263 |
| 2 | Oct 10, 2017 | 629.55 | 3.1480 | 631.514648 | -1.964648 |
| 3 | Oct 11, 2017 | 626.75 | 3.1775 | 627.256547 | -0.506547 |
| 4 | Oct 12, 2017 | 624.47 | 3.2020 | 623.720158 | 0.749842 |

```
In [8]: # Plot of the scattered data, the OLS regression, and the sum of squared errors

plt.figure(figsize=(15, 15))
plt.title('The sum of squared errors is {}'.format(error_before))
plt.scatter(df_before['Copper_price'], df_before['USD_CLP'], color='darkblue', label='Data, Oct 06, 2017 - Oct 04 2019')
plt.plot(df_before['Copper_price'], df_before['Best_fit'], color='red',
label='OLS fit, Oct 06, 2017 - Oct 04 2019')
plt.xlim(df_before['Copper_price'].min()-0.1, df_before['Copper_price'].max()+0.1)
plt.ylim(df_before['USD_CLP'].max()+20, df_before['USD_CLP'].min()-20)
plt.xlabel("Copper price")
plt.ylabel("USD-CLP")
plt.legend()
plt.show()
```

The sum of squared errors is 92779.8



```
In [21]: # We give with .summary() a summary of the performance of the model before the civil unrest.  
# Contains the p-value of 'Copper_price', the confidence interval and the coefficient of determination  
  
fit_before.summary()
```

Out[21]: OLS Regression Results

| | | | |
|--------------------------|------------------|----------------------------|-----------|
| Dep. Variable: | USD_CLP | R-squared: | 0.855 |
| Model: | OLS | Adj. R-squared: | 0.854 |
| Method: | Least Squares | F-statistic: | 2996. |
| Date: | Wed, 23 Feb 2022 | Prob (F-statistic): | 1.22e-215 |
| Time: | 14:06:10 | Log-Likelihood: | -2057.6 |
| No. Observations: | 512 | AIC: | 4119. |
| Df Residuals: | 510 | BIC: | 4128. |
| Df Model: | 1 | | |
| Covariance Type: | nonrobust | | |

| | coef | std err | t | P> t | [0.025 | 0.975] |
|---------------------|-----------|---------|---------|-------|----------|----------|
| Intercept | 1085.9046 | 7.831 | 138.663 | 0.000 | 1070.519 | 1101.290 |
| Copper_price | -144.3424 | 2.637 | -54.735 | 0.000 | -149.523 | -139.162 |

| | | | |
|-----------------------|--------|--------------------------|--------|
| Omnibus: | 8.791 | Durbin-Watson: | 0.132 |
| Prob(Omnibus): | 0.012 | Jarque-Bera (JB): | 5.680 |
| Skew: | -0.091 | Prob(JB): | 0.0584 |
| Kurtosis: | 2.517 | Cond. No. | 43.4 |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

We do now the same analysis for the observed data after the civil unrest, during the time frame Oct 08, 2019 - Oct 08 2021.

```
In [23]: # Store the csv files in dataframes pesos_after and copper_after, and order them by ascending date

pesos_after = pd.read_csv('CLP_after.csv')
copper_after = pd.read_csv('Copper_after.csv')
pesos_after, copper_after = pesos_after[:::-1], copper_after[:::-1]
```

```
In [24]: # Merge the data into one dataframe 'df_after', with 'Date' as (inner) merge criterion

df_after = pd.merge(left=pesos_after, right=copper_after, left_on='Date', right_on='Date')
df_after = [df_after['Date'], df_after['Price_x'], df_after['Price_y']]
headers = ['Date', 'USD_CLP', 'Copper_price']
df_after = pd.concat(df_after, axis=1, keys=headers)
```

```
In [25]: # The attribute 'params' returns the list of estimated parameters for the model.
# These are the intercept and the slope of the OLS regression, params[0] and params[1], respectively.

fit_after = smf.ols(formula='USD_CLP~Copper_price', data=df_after).fit()
intercept_after, slope_after = fit_after.params[0], fit_after.params[1]
```


```
In [26]: # Equation of OLS regression (with column 'Best_fit'), and error between the 'Best_fit' and the obser

df_after['Best_fit'] = slope_after*df_after['Copper_price'] + intercept_after
df_after['Error'] = df_after['USD_CLP'] - df_after['Best_fit']
error_after = round((df_after['Error']**2).sum(), 1)
print(df_after.head())
```

| | Date | USD_CLP | Copper_price | Best_fit | Error |
|---|--------------|---------|--------------|------------|------------|
| 0 | Oct 08, 2019 | 725.24 | 2.6350 | 793.044021 | -67.804021 |
| 1 | Oct 09, 2019 | 723.20 | 2.6315 | 793.168131 | -69.968131 |
| 2 | Oct 10, 2019 | 718.48 | 2.6720 | 791.732003 | -73.252003 |
| 3 | Oct 11, 2019 | 709.48 | 2.6890 | 791.129184 | -81.649184 |
| 4 | Oct 14, 2019 | 712.43 | 2.6950 | 790.916425 | -78.486425 |

In [27]: *# Plot of the scattered data, the OLS regression, and the sum of squared errors*

```
plt.figure(figsize=(15, 15))
plt.title('The sum of squared errors is {}'.format(error_after))
plt.scatter(df_after['Copper_price'], df_after['USD_CLP'], color='darkblue', label='Data, Oct 08, 201')
plt.plot(df_after['Copper_price'], df_after['Best_fit'], color='red',
label='OLS fit, Oct 08, 2019 - Oct 08 2021')
plt.xlim(df_after['Copper_price'].min()-0.1, df_after['Copper_price'].max()+0.1)
plt.ylim(df_after['USD_CLP'].max()+20, df_after['USD_CLP'].min()-20)
plt.xlabel("Copper price")
plt.ylabel("USD-CLP")
plt.legend()
plt.show()
```



In [28]: *# We give with .summary() a summary of the performance of the model after the civil unrest.*

```
# Contains the p-value of 'Copper_price', the confidence interval and the coefficient of determination
fit_after.summary()
```

Out[28]:

| OLS Regression Results | | | | | | |
|------------------------|------------------|-------------------|---------|---------------------|----------|---------|
| Dep. Variable: | USD_CLP | | | R-squared: | 0.409 | |
| Model: | OLS | | | Adj. R-squared: | 0.407 | |
| Method: | Least Squares | | | F-statistic: | 359.2 | |
| Date: | Wed, 23 Feb 2022 | | | Prob (F-statistic): | 2.70e-61 | |
| Time: | 14:09:59 | | | Log-Likelihood: | -2538.6 | |
| No. Observations: | 522 | | | AIC: | 5081. | |
| Df Residuals: | 520 | | | BIC: | 5090. | |
| Df Model: | 1 | | | | | |
| Covariance Type: | nonrobust | | | | | |
| | coef | std err | t | P> t | [0.025 | 0.975] |
| Intercept | 886.4810 | 6.412 | 138.252 | 0.000 | 873.884 | 899.078 |
| Copper_price | -35.4599 | 1.871 | -18.953 | 0.000 | -39.135 | -31.784 |
| Omnibus: | 2.696 | Durbin-Watson: | 0.042 | | | |
| Prob(Omnibus): | 0.260 | Jarque-Bera (JB): | 2.693 | | | |
| Skew: | 0.138 | Prob(JB): | 0.260 | | | |
| Kurtosis: | 2.782 | Cond. No. | 17.3 | | | |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.